


GigaVoxels: Ray-Guided Streaming for Efficient and Detailed Voxel Rendering

Cyril Crassin, Fabrice Neyret, Sylvain Lefebvre, Elmar Eisemann

summarized by Sarah Kushner

Volume Rendering and Voxels

Used to represent pseudo-surfaces like foliage, clouds, smoke and even extremely detailed geometry



These two techniques together are used in conjunction to deal with the level-of-detail problem.



Proposal

Current hardware is ready to render massive amounts (several billion voxels) of volume data in real-time or at an interactive rate.

Problems

1. Memory limitations

- a. The transfer of 512MB of data onto the GPU prevents real-time performance.

2. Costly rendering

- a. Costly because of large amounts of voxels.

Approach


- The norm for concentration of scene details in CG is usually an interface between dense clusters of detail and empty space.
- Focuses on out-of-core voxel rendering.

Optimizations

- Depending on the camera, not everything needs to be in memory all the time.
- Distant objects can have lower mipmap levels, less detail, and in turn take less memory.

Data Structures

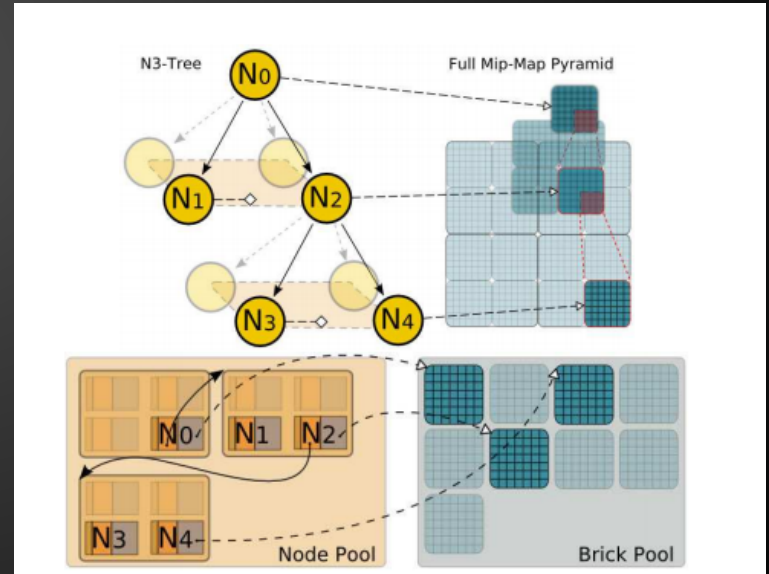
- N^3 trees and mipmapped texture tiles.
- Each node is single pointer to a "brick" or an empty space.




- Brick: small voxel grid of size M^3 (usually $M = 32$) that approximates part of the original volume.
- All bricks stored in brick pool.
- Updates only triggered if data in the brick pool is missing.
- New bricks stored in oldest locations according to timestamp. (Last Recently Used)

Organization

- All nodes in 3D texture called "node pool".
- Grouping gains access to all children through 1 pointer.
- 16x memory improvement since nodes are now 64 partitioned accordingly.



Nodes



- 30 bits - encode a pointer to the child nodes
- 1 bit - whether the node is refined to a maximum or contains more data
- 1 bit - whether the content is a constant RGBA8 value or described by a brick
- 30/32 bits - either a pointer to an M^3 brick or the average value at this location

Rendering

- March the data in the structure along the view rays while accumulating color and opacity.
- Iterative descent from the tree root like the kd-restart algorithm.
- Level-of-detail determined in shader.
- Ray-casting in one big fragment shader.
- Adapt sampling rate & mipmap level during the ray marching depending on viewpoint.

Communication with CPU


- Brick pool on CPU.
- Texture-update calls -> modifications go to GPU (Node pool).
- Unified management of the brick pool and the node pool as two Last Recently Used controlled caches.

Node List

- During rendering, traverse the node tree. Stop on node corresponding to the needed level-of-detail.
- If the required data is present, we traverse the mipmapped brick.
- For CPU to keep correct nodes in pool, collect the current node index in a node-list.
- If level-of-detail is missing or the node is terminal, it is added to the node-list.

Neighbor Rays

- Group rays in packages in the image plane.





- Sort to ensure only one instance of each node reference.

- Sorting can also be costly.

- Instead of brute-force sorting, rays are only compared with neighboring rays.

Results

- 20 octaves of Perlin noise, shading, materials
- 20 fps for a 1024^3 volume




- 1024^3 scanned volume of bone
- $N = 2$ and $M = 16$
- 60 fps

Conclusion

Through compact data structures and limited, clever CPU to GPU communication, interactive and real-time rendering speeds are possible for extremely large volume data.

Demo



<https://www.youtube.com/watch?v=HScYuRhgEJw>